

# 6SENG002W Concurrent Programming

## Lecture 11

### *Translating FSP Models into Concurrent Java Programs*



# Translating FSP Models into Concurrent Java Programs

The aim of this lecture is to illustrate how to translate FSP into Java.

We shall do this by considering:

- ▶ The following FSP program examples:
  - ▶ the *Ornamental Garden Problem* – involving processes interacting & sharing an object using *mutual exclusion*.
  - ▶ a *Semaphore* – involving *conditional synchronisation*.
- ▶ The issues related to Threads *interacting* & *sharing objects using Mutual Exclusion*.
- ▶ How to translate various aspects of FSP processes & programs into Java threads, monitors & programs:
  - ▶ an *FSP process* into either a Java *monitor* or *thread*,
  - ▶ an *FSP composite process* (program) into a Java *main program*.

**Note:** these lecture notes are based on material in Chapters 4, 5, & 7 of the recommended book: *Concurrency: State Models & Java Programs, (2nd Edition)*, J. Magee & J. Kramer, Wiley, 2006. (ISBN 978-0-470-09355-9)

## PART I

*The Ornamental Garden Problem  
an Example of  
Process Interaction & Mutual Exclusion*

## The Ornamental Garden Problem

To illustrate the issues of *thread interaction & mutual exclusion*, we use the *Ornamental Garden problem*. (See the recommended book Chapter 4.)




Figure : 11.1 Ornamental Garden

- ▶ People visit an ornamental garden, by entering through either of *two turnstiles* – East & West.
- ▶ To simplify the problem people are allowed to *enter* but *never leave*.
- ▶ Management needs to determine how many people are in the garden at any one time, so a single *people counter* is required.
- ▶ So when a person passes through either turnstile, the turnstile must increment the people counter, therefore the counter must be *shared*.
- ▶ The FSP program used to model this system must represent the:
  - ▶ *shared people counter*, that requires mutually exclusive access,
  - ▶ *East & West turnstiles* that have *shared ME access* to the people counter.

## FSP Model of the Ornamental Garden

In general *each object* or *set of objects* will be **modelled** as an *FSP process*.

The Ornamental Garden FSP program models the *increment* action by incrementing a variable counter process `SAFE_COUNTER`, that describes the read & write accesses to a variable.

There is **no explicit mention** of an *increment* action, instead, increment is modelled using `read` & `write` actions by `TURNSTILE`'s local process `INCREMENT`.

The **east** & **west** `TURNSTILE` processes, each has its own copy of the `read` & `write` actions that make up the increment operation.

In other words, the increment operation is **not an atomic action**.

Therefore, *interference*, i.e. destructive updates, can occur by means of the arbitrary interleaving of `read` & `write` actions.

The *solution to interference* is to give operations that access a shared object *mutually exclusive* access to that object.

This ensures that an update, i.e. an increment (`read` then `write`), is **not interrupted** by concurrent updates.

## Ornamental Garden: Preliminary Definitions

We use the following preliminary definitions of: a *constant*, a *number range*, an *action set* & 2 *process label sets*.

These will be used in the definition of the FSP processes that model the Ornamental Garden problem.

### Ornamental Garden: Definitions

```
const  N = 4           /* Max value for counter */

range CR = 0..N        /* Range for counter */

/* Alphabet of SAFE_COUNTER process */
set CounterAlpha = { counter.{ read[CR], write[CR],
                             acquire, release } }

/* Process labels for the active processes */
set Turnstiles   = { east, west }

set CounterUsers = { east, west, management }
```

**CounterAlpha** is the counter process's alphabet & is used to extend the alphabets of the two turnstiles **east** & **west**.

## Ornamental Garden: an “Unsafe” Shareable Counter

We begin by defining an FSP process `UNSAFE_COUNTER` that models a *counter variable* & the *basic operations* on the *counter*.

This process **DOES NOT** attempt to enforce *mutual exclusion*, this will be added in the next stage.

```
Ornamental Garden: UNSAFE_COUNTER

UNSAFE_COUNTER      = US_COUNT[ 0 ] ,

US_COUNT[ v : CR ] = (  read[ v ]          -> US_COUNT[ v ]
                       | write[ nv : CR ] -> US_COUNT[ nv ] ) .
```

Where:

- ▶ the counter is *initialised* to 0 by `US_COUNT[ 0 ]`,
- ▶ the `read[ v ]` action is used to *read/get* the current value of the counter `v`,
- ▶ the `write[ nv : CR ]` action is used to *write/set* the new value of the counter to `nv`.

## Ornamental Garden: Expanded UNSAFE\_COUNTER

To understand how this process UNSAFE\_COUNTER models a counter here is the expanded version:

Expanded: UNSAFE\_COUNTER

```
UNSAFE_COUNTER      = US_COUNT[ 0 ] ,

US_COUNT[ 0 ] = (  read[ 0 ]          -> US_COUNT[ 0 ]
                  | write[ nv : CR ] -> US_COUNT[ nv ] ) ,

US_COUNT[ 1 ] = (  read[ 1 ]          -> US_COUNT[ 1 ]
                  | write[ nv : CR ] -> US_COUNT[ nv ] ) ,

US_COUNT[ 2 ] = (  read[ 2 ]          -> US_COUNT[ 2 ]
                  | write[ nv : CR ] -> US_COUNT[ nv ] ) ,

US_COUNT[ 3 ] = (  read[ 3 ]          -> US_COUNT[ 3 ]
                  | write[ nv : CR ] -> US_COUNT[ nv ] ) ,

US_COUNT[ 4 ] = (  read[ 4 ]          -> US_COUNT[ 4 ]
                  | write[ nv : CR ] -> US_COUNT[ nv ] ) .
```

So the local process `US_COUNT[0]` represents a *counter* with the value 0, similarly for the others.



## Ornamental Garden: Expanded UNSAFE\_COUNTER's US\_COUNT [ 0 ]

We can further expand UNSAFE\_COUNTER by expanding its local process US\_COUNT [ 0 ], the expanded version:

```
Expanded: US_COUNT [ 0 ]  
  
US_COUNT [ 0 ] = (  read [ 0 ]  -> US_COUNT [ 0 ]  
                   | write [ 0 ] -> US_COUNT [ 0 ]  
                   | write [ 1 ] -> US_COUNT [ 1 ]  
                   | write [ 2 ] -> US_COUNT [ 2 ]  
                   | write [ 3 ] -> US_COUNT [ 3 ]  
                   | write [ 4 ] -> US_COUNT [ 4 ] )
```

US\_COUNT [ 0 ] represents a *counter* with the value 0, its value can be read & it can be set to any value 0 – 4 by write.

The other local processes: US\_COUNT [ 1 ] – US\_COUNT [ 4 ], can be expanded similarly.

## Ornamental Garden: a “Safe” Shareable Counter

To ensure that the `UNSAFE_COUNTER` counter can be “*safely shared*” access to it by the 2 turnstiles must be *mutually exclusive*.

This is achieved by combining it with a *locking* process.

A *lock* can be modelled by the `LOCK` process & then combined with the `UNSAFE_COUNTER` process by the composition `SAFE_COUNTER`:

```
Ornamental Garden: VISITORS_COUNTER

LOCK = ( acquire -> release -> LOCK ) .

|| SAFE_COUNTER = ( LOCK || UNSAFE_COUNTER ) .

|| VISITORS_COUNTER = ( counter : SAFE_COUNTER ) .
```

We use the “*safe*” `VISITORS_COUNTER` version in the complete system, rather than the “*unprotected*” & “*unsafe*” `UNSAFE_COUNTER` version.

This is just `SAFE_COUNTER` labelled with “*counter*”, which in effect collects the actions together & defines a “*counter interface*”.

# Ornamental Garden: Defining a TURNSTILE Process

Finally, the definition of the TURNSTILE process involves three steps.

## Step 1: Communication between a TURNSTILE & the VISITORS\_COUNTER

The 2 processes communicate using *inter-process communication* (IPC), by *synchronising* the two actions `read[v]` & `write[nv]`:

1. The TURNSTILE process “*reads*” the counter’s value from the VISITORS\_COUNTER by them performing a *synchronised* `read[v]` action:
  - ▶ VISITORS\_COUNTER *outputs* `v`, the *current value of the counter*
  - ▶ TURNSTILE *inputs* the value of `v`.
2. The TURNSTILE process “*writes*” the counter’s new value to the VISITORS\_COUNTER by them performing a *synchronised* `write[nv]` action:
  - ▶ TURNSTILE *outputs* `nv`, the *new value of the counter*,
  - ▶ VISITORS\_COUNTER *inputs* `nv`.

## Defining a TURNSTILE Process: Step 2

**Step 2:** TURNSTILE uses Mutual Exclusion Protocol to Access the VISITORS\_COUNTER's value

Each TURNSTILE process **must be required to adhere** to the *mutual exclusion protocol* (MEP):

MEP = **lock** -> ``use resource`` -> **unlock**

We implement this in FSP by requiring the TURNSTILE process to:

1. **acquire** the lock before accessing the **counter** variable,
2. update the the **counter** variable, by *synchronising* the two actions `read[v]` & `write[nv]`,
3. **release** the lock after it has finished updating the **counter** variable.

## Defining a TURNSTILE Process: Step 3

### Step 3: Block Resource Performing Asynchronous Actions

Ensuring Step 2 implements ME correctly, we need to make sure that the VISITORS\_COUNTER **cannot do any unauthorised asynchronous actions**.

```
alphabet ( VISITORS_COUNTER ) = CounterAlpha  
  
= { counter.{ read[CR], write[CR], acquire, release } }  
  
= { counter.{ read[0], read[1], read[2], read[3], read[4],  
              write[0], write[1], write[2], write[3], write[4],  
              acquire, release } }
```

So the actions that would cause a problem by being performed *asynchronous* by VISITORS\_COUNTER are: write[0] – write[4].

This is because a turnstile will **only every attempt to write one of the values 0 – 4 at a time**, but VISITORS\_COUNTER is **always willing to do all of them**. (See expanded US\_COUNT[0] above.)

So VISITORS\_COUNTER **must be stopped** from doing a write *asynchronously* that a turnstile **does not** want to perform.

This is prevented by requiring all actions of VISITORS\_COUNTER to be *synchronised* with the 2 turnstile processes.

## Ornamental Garden: TURNSTILE Processes

First define the single TURNSTILE process.

```
Ornamental Garden: TURNSTILE
TURNSTILE = ( open -> OPEN_TURNSTILE ) ,

OPEN_TURNSTILE = (  arrive -> INCREMENT
                    | close -> END ) ,

INCREMENT
= (  counter.acquire      -> /* Lock counter */
    counter.read[ x : CR ] -> /* Get counter value */
    counter.write[ x + 1 ] -> /* Write new counter value */
    counter.release       -> /* Unlock counter */
    OPEN_TURNSTILE
  ) +CounterAlpha .
```

Requiring **all actions** of VISITORS\_COUNTER to be *synchronised* with a turnstile process is achieved by *extending the alphabet* of TURNSTILE's local process INCREMENT with the alphabet of VISITORS\_COUNTER using **CounterAlpha**.

The result is that VISITORS\_COUNTER can only do an action if the TURNSTILE is willing to do it as well, i.e. *synchronise* on it.

## Ornamental Garden: East & West TURNSTILE Processes

We can now use this to define the two *East* & *West* TURNSTILE Processes:

Ornamental Garden: TURNSTILES

```
|| EAST_TURNSTILE = ( east : TURNSTILE ) .  
  
|| WEST_TURNSTILE = ( west : TURNSTILE ) .  
  
|| TURNSTILES = ( WEST_TURNSTILE || EAST_TURNSTILE )  
                  / { open / Turnstiles.open,  
                    close / Turnstiles.close } .
```

```
alphabet( EAST_TURNSTILE )  
= { east.{ open, arrive, close,  
           counter.acquire, counter.release,  
           counter.read[CR], counter.write[CR] } }  
  
= { east.open, east.arrive, east.close,  
    east.counter.acquire, east.counter.release,  
    east.counter.read[CR], east.counter.write[CR] }
```

Note that “open / Turnstiles.open” means relabel both *east*.open & *west*.open to open, similarly for close.

## Ornamental Garden: MANAGEMENT process

The MANAGEMENT process can check the value of the counter at any time using the DISPLAY process to read the value of the **counter**, without having to lock & unlock it.

In this case this is acceptable since the DISPLAY process does not update the value of the **counter** variable.

The close action closes the garden & terminates the system.

```
Ornamental Garden: MANAGEMENT

DISPLAY = (   counter.read[CR] -> DISPLAY
            | close                -> END      ) +CounterAlpha .

|| MANAGEMENT = ( management : DISPLAY )
                  / { close / management.close } .
```

As with INCREMENT, its alphabet must be *extended* with that of the VISITORS\_COUNTER's to ensure **no inappropriate actions** are performed by the VISITORS\_COUNTER.

For example, any of the **counter**.write[CR] ones, e.g. **counter**.write[0], **counter**.write[1], etc.



## Ornamental Garden: Complete FSP Model

The complete Ornamental Garden system is defined by the GARDEN process:

Ornamental Garden: GARDEN

```
|| GARDEN = (
    || TURNSTILES
    || MANAGEMENT
    || CounterUsers :: VISITORS_COUNTER
) .
```

which expands to:

Expanded GARDEN

```
|| GARDEN =
(
    (west : TURNSTILE) /{open/west.open, close/west.close}
    ||
    (east : TURNSTILE) /{open/east.open, close/east.close}
    ||
    ( management : DISPLAY )
    ||
    ( { east, west, management } :: counter : LOCK )
    ||
    ( { east, west, management } :: counter : SAFE_COUNTER )
) .
```

## Notes on the Ornamental Garden FSP Model

The alphabet for the `TURNSTILE` process is extended with the **CounterAlpha** set using the alphabet extension construct  $+ \{ \dots \}$ .

This is to ensure that **no unintended asynchronous actions** are performed within the system.

For example, if a `VISITORS_COUNTER` `write[nv]` is **not shared** (synchronised) with another process, i.e. a `TURNSTILE` or `DISPLAY` then it can do it **asynchronous**.

`TURNSTILE` **never** does the action “`counter.write[0]`” since it *always increments the value* it reads.

But, since “`counter.write[0]`” is in **CounterAlpha**, `VISITORS_COUNTER` is **prevented from preforming it asynchronously**.

The `close` action is synchronised on by the turnstiles & the management processes it terminates the system, it is only accepted as an alternative to an `arrive` action.

Note that `VISITORS_COUNTER` is **shared by** the **east** & **west** `TURNSTILES` & `MANAGEMENT`.

## Ornamental Garden: Structure Diagram

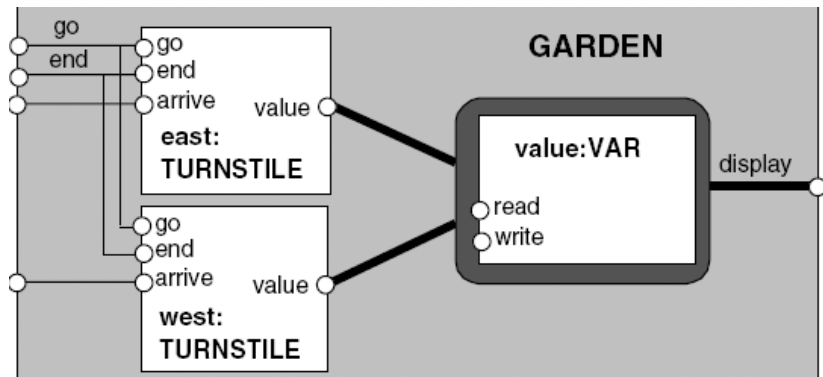


Figure : 11.2 Ornamental Garden Structure Diagram

Notes: in the diagram “value” is **counter** & “value:VAR” is VISITORS\_COUNTER, i.e. **counter**:SAFE\_COUNTER.

## PART II

# *Threads Communicating/Interacting by Sharing an Object using Mutual Exclusion*

## Threads Communicating via a *Shared Object*

One of the simplest ways for two or more Java threads to *"communicate"* &/or *"interact"* with each other is via a *shared object*.

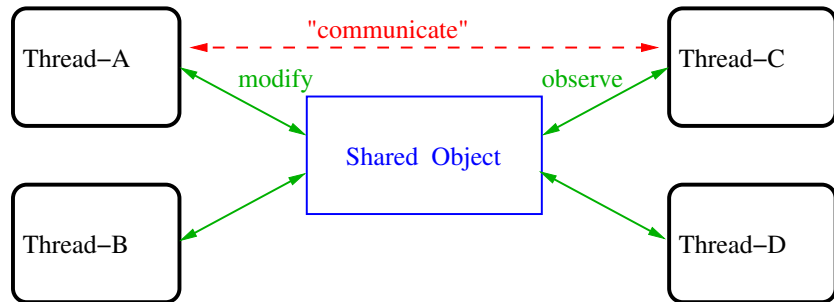


Figure : 11.3 Threads Communicate via a Shared Object

## Mutual Exclusion: Context – *Shared Objects*

The *shared object* from Fig. 11.3, that is intended to be used in this way, must have an *interface* that allows this.

That is it must have `public` methods that allow its *state* to be “*modified*” & “*observed*” by these methods.

The collection of threads can then communicate & interact by invoking the shared object’s `public` methods.

Consequently, two threads can **communicate** by one thread *modifying* the state of the shared object & the other thread then *observing* the state.

For example, Thread–A & Thread–C.

Other examples of threads sharing an object, are when a collection of threads *cooperate to update* information encapsulated in a shared object.

## Mutual Exclusion: Problem – *Interference*

We have previously seen that the execution of the instructions from a collection of processes/threads can be interleaved in an arbitrary fashion.

This interleaving can result in incorrect updates to the state of a shared object, resulting in **corrupted data**, this is known as **interference**.

These types of objects are sometimes referred to as *damaged objects*.

This is clearly a situation that must be avoided at all costs.

## Mutual Exclusion: Solution – *Java Monitors*

The problem of **interference**, is solved by the enforcement of *mutually exclusive* access to the shared object, by means of a *locking* arrangement.

In Java this means ensuring that the *shared object* is implemented as a **correctly functioning “secure” Java monitor**.

That is, the Java monitor ensures:

- ▶ *complete data encapsulation inside the monitor*
  - **all data** is either `private` or `protected`
- ▶ *execution of all its public methods are mutually exclusive*
  - **all public methods** are `synchronized`.



## PART III

### *Translating FSP into Java*

## Translating FSP into Java

An FSP program, i.e. a collection of FSP processes, is used to model a program.

In a given FSP model of a system each:

**Shared object:** is modelled as an *FSP process*,

**Process:** is also modelled as an *FSP process*.

However, FSP processes (that make up a program), do not distinguish between:

- ▶ *“Passive” entities* (shared objects: Java monitors).
- ▶ *“Active” entities* (processes: Java threads).

They are both modelled as finite state machines.

This uniform treatment facilitates analysis.

Therefore, in translating an FSP program into a Java program *we must decide* what kind of Java *“object”* each individual FSP process is modelling.

## Mutual Exclusion in Java

In Java usually represent “*shared objects*” as *Java monitors*.

At the Java monitor level of abstraction, we can ignore the details of *locks* & *mutual exclusion* (as provided by the use of `synchronized` methods).

Concurrent calls of a Java `synchronized` method are *mutually exclusive*.

So we can translate the FSP `SAFE_COUNTER` process into a Java monitor & make the `increment` method **synchronized**.

```
class SAFE_COUNTER          // a ``monitor``
{
    private int value = 0 ;

    SAFE_COUNTER( int initial ){ value = initial ; }

    public synchronized void increment()
    {
        int temp = value ;      // "read[ x : CR ]"
        value     = temp + 1 ;  // "write[ x + 1 ]"
    }
}

SAFE_COUNTER counter = new SAFE_COUNTER( 0 ) ; // "US_COUNT[ 0 ]"
```

## Java Monitor Semantics Ensures Mutual Exclusion

Java associates a *lock* with every object.

The Java compiler inserts code to *acquire the lock* before executing the body of a `synchronized` method & code to *release the lock* before the method returns.

Concurrent threads *trying to access the monitor* are **blocked** until the lock is *released*.

Since only one thread at a time may hold the lock, only one thread may be executing the `synchronized` method.

If this is the only method, as in the example, *mutual exclusion* to the shared object is ensured.

If an object has *more than one method*, to ensure mutually exclusive access to the state of the object, **all the methods should be synchronized**.

Access to an object may also be made *mutually exclusive* by using the `synchronized` statement, but this is **not recommended**, as it is too *"unstructured"* & prone to errors.

## Summary

We have discussed *thread interaction* via *shared objects*.

The Ornamental Garden example served to demonstrate that uncontrolled interleaving of method instructions leads to **interference**, i.e. **destructive update** of the state of the shared object.

**Interference** can be avoided by giving each concurrent method activation *mutually exclusive access to the shared state*.

In Java, this is achieved by making such methods `synchronized`.

`synchronized` methods **acquire a lock** associated with the object before accessing the object state & **release the lock** after access.

Since only one thread at a time can acquire the lock, `synchronized` methods obtain *mutually exclusive access* to the object state.

The answer is to: **ensure that all the methods of objects shared between threads are synchronized**.

They can then be treated as *“atomic actions”* for modelling purposes.

Note that **interference bugs** in real concurrent programs are notoriously difficult to find.

## Condition Synchronisation: Semaphores

We illustrate *condition synchronisation* using a simple version of Dijkstra's semaphores as an example.

Recall that a semaphore  $s$  is an integer variable that can take only non-negative values.

Once  $s$  has been given an initial value, the only operations permitted on  $s$  are `s.claim()` & `s.release()` defined as follows:

```
s.claim(): when ( s > 0 ) { decrement s ; }  
          when ( s == 0 ) { wait/suspend ; }  
  
s.release(): increment s ;
```

In the following, we describe how *semaphores* can be modelled by an *FSP process* & how that can be implemented using a *Java monitor*.

However, in practice, semaphores are a low-level mechanism & are sometimes used to implement the higher-level monitor construct, rather than vice versa.

# FSP Semaphore Model & Process

First step in modelling a system: decide which *events* or *actions* are of interest.

For a semaphore there are only two actions:

- ▶ **claim**
- ▶ **release.**

Second step: identify the *processes* – the *users* of the shared variable.

The control requirements for a semaphore's `claim` action are modelled using the FSP *guarded action construct*:

```
when ( B ) action
```

## FSP Semaphore: SEMAPHORE

The FSP process to model this simple view of a semaphore is:

```

                                SEMAPHORE
const MAX_NAT = 3
range NAT      = 0..MAX_NAT

SEMAPHORE( N = 0 ) = SEMA[ N ] ,

SEMA[ s : NAT ] = (  when ( s > 0 ) claim -> SEMA[ s - 1 ]
                    |  release          -> SEMA[ s + 1 ] ),

SEMA[ MAX_NAT + 1 ] = ERROR .
```

FSP can **only model finite concurrent systems**, so we only model semaphores that take a *finite range of values*, i.e. 0 – 3.

SEMAPHORE allows a maximum of three **claim** actions to be accepted before a **release** action must occur.

If a **release** action results in a semaphore value of 4, then it is treated as an error & is defined as the error process ERROR.



## Behaviour of the FSP Semaphore Process

The semaphore's behaviour is illustrated in the LTSA diagram Fig. 11.4.




Figure : 11.4 SEMAPHORE LTS

This directly models the first (simple) definition for a semaphore given above.

**claim** is only accepted when ( $s > 0$ ) in the SEMAPHORE.

**release** is **not guarded** (or the guard is just `true`).

SEMAPHORE has values in the range  $0 \dots \text{MAX\_NAT}$  & has an initial value `N`.

If a **release** causes `MAX_NAT` to be exceeded then SEMAPHORE moves to the ERROR state **-1**.

When SEMAPHORE is used in a larger model, we must ensure that this ERROR state **does not occur**.

## FSP Semaphore Example

We use a semaphore `mutex`, shared by three processes `mp[1..3]`, to ensure *mutually exclusive* access to some resource.

Each process performs the action `mutex.claim` to get exclusive access & `mutex.release` to release it.

Access to the resource (via critical section), is modelled by the action `critical[i]`.

The composite process `SEMADEMO`, which combines the three processes `mp[i] : MutexProcess` & a semaphore is:

```
SEMADEMO
range IDS = 1..3
MutexProcess( ID = 0 ) = ( mutex.claim ->
                           critical[ID] ->
                           mutex.release -> MutexProcess ) .

|| MutexProcesses = ( forall [ i : IDS ]
                      ( mp[ i ] : MutexProcess(i) ) ) .

|| SEMADEMO = (      MutexProcesses
                 || { mp[IDS] } :: mutex : SEMAPHORE(1) ) .
```

# Ensuring Mutual Exclusion Using the Semaphore

To achieve *mutual exclusion* use a *binary semaphore* initially **unlocked**, i.e. 1.

The first process that tries to execute its critical action, performs a **mutex.claim** making the value of mutex zero.

No further process can perform **mutex.claim** until the original process releases mutual exclusion by **mutex.release**.

## Condition Synchronisation in Java (Monitors)

Java *monitor objects* have a *synchronisation lock* & a thread *wait set*.

Recall the following methods provided by the `Object` class:

- ▶ `wait()` – place current thread in the wait set.
- ▶ `notify()` – wake up a single thread from wait set.
- ▶ `notifyAll()` – wake up all threads in wait set.

The operations **fail** if called by a thread that does **not currently “own” the monitor**, i.e. has the synchronisation lock.

A thread enters a monitor when it **acquires the mutual exclusion lock** associated with the monitor & exits the monitor when it **releases the lock**.

A thread calling `wait()` *exits the monitor*.

This allows other threads to enter the monitor &, when the appropriate condition is satisfied, to call `notify()` or `notifyAll()` to awake waiting threads.

## Condition Synchronisation in FSP

The basic format for modelling a *guarded action* for some condition **cond** & action **action** using FSP is shown below:

```
FSP:    when ( cond ) action -> NEWSTATE
```

The corresponding format for implementing the *guarded action* for condition **cond** & action **action** using Java is as follows:

```
public synchronized void action( )  
    throws InterruptedException  
{  
    while ( !cond )    // "when ( cond )"  
    {  
        wait() ;        // cond = false  
    }  
    // cond = true  
  
    // "action -> NEWSTATE"  
    // modify monitor data "action -> NEWSTATE"  
  
    notifyAll() ;  
}
```

## Notes on the Translation

The `while` loop is necessary to ensure that **cond** is indeed satisfied when a *thread re-enters the monitor*.

Although the thread invoking `wait()` may have been notified that **cond** is satisfied, thereby releasing it from the monitor wait set.

But **cond** may be **invalidated by another thread** that runs between the time that the waiting thread is awakened & the time it re-enters the monitor (by acquiring the lock).

If an action *modifies the data of the monitor*, it can call `notifyAll()` to awaken all other threads that may be waiting for a particular condition to hold with respect to this data.

If it is not certain that only a single thread needs to be awakened, it is safer to call `notifyAll()` than `notify()` to make sure that threads are not kept waiting unnecessarily.

## Translation General Rules

**General Rules:** for guiding the translation of an *FSP process model* into a *Java monitor* are as follows:

1. In general if an FSP process that represents a shared object offers a *choice* (“|”) between several *unguarded actions* & *guarded actions*, then as a first attempt at translation each one should be implemented as a separate `synchronized` method.
2. The *choice* (“|”) is represented in the monitor by the fact that a user of the shared resource (monitor) is “*offered the choice*” of calling any one of the `synchronized` methods.
3. Each *guarded action* in the FSP model of a monitor is implemented as a:
  - ▶ `synchronized` method
  - ▶ that uses a `while` loop & `wait()` to implement the guard.
4. The `while` loop **condition** is the **logical negation** (!) of the **model guard condition**.

## Translating the FSP Semaphore into Java

There is of course Java's pre-defined semaphore class:

```
java.util.concurrent.Semaphore
```

that we briefly looked at in a previous lecture.

But we shall provide our own explicitly defined Java Semaphore class.

Semaphores are *passive* objects that react to `claim` & `release` actions; they do **not initiate actions**.

Consequently, we implement a semaphore as a Java *monitor* class.

The actions `claim` & `release` become `synchronized` methods of our monitor semaphore class.



## Translating the FSP Semaphore into Java: Guard

The guard: “**when** (  $v > 0$  )”, on the `claim` action in the FSP process model, is implemented using *conditional synchronisation*.

That is, we must use the *monitor guard* `while` loop:

```
while ( BExp )  
{  
    wait() ;  
}
```

The `claim` changes the state of the monitor by decrementing its value `s`.

However, we do **not use** `notify()` to *signal the change in state*.

This is because threads only wait for the value of the semaphore to be *incremented*, they **do not wait** for the value to be *decremented*.

## Implementing FSP's SEMAPHORE Process using a Java Semaphore class

Note that this is a very simple implementation of a semaphore.

```
public class Semaphore
{
    private int value ;

    public Semaphore ( int initial ) { value = initial ; }

    public synchronized void claim( ) throws InterruptedException
    {
        while ( value == 0 )
        {
            wait() ;
        }
        value-- ;
    }

    public synchronized void release( )
    {
        value++ ;
        notify() ;
    }
}
```

## Implementing FSP's `MutexProcess` using a Java Thread

The FSP `MutexProcess` process is implemented as a Java Thread class.

Each one performs the action `mutex.claim()` to get exclusive access & `mutex.release()` to release it.

```
class MutexProcess extends Thread
{
    private Semaphore mutex ;

    MutexProcess( Semaphore sema )
    {
        super( "MutexProcess" ) ; // Thread constructor
        mutex = sema ;
    }

    public void run()
    {
        try { while ( true )
            {
                mutex.claim() ;           // get ME
                // critical section
                mutex.release() ;         // release ME
            }
        } catch (InterruptedException e){}
    }
} // MutexProcess
```

## Implementing FSP's SEMADEMO Process using a Java Thread class

The Java version of the composite process SEMADEMO, which combines the three processes MutexProcess & a semaphore is:

```
class SEMADEMO
{
    public static void main( String args[] )
    {
        final int UNLOCKED = 1 ;

        Semaphore mutex = new Semaphore( UNLOCKED ) ;

        Thread mp[] = new Thread[3] ;

        // create 3 MutexProcess
        mp[0] = new MutexProcess( mutex ) ;
        mp[1] = new MutexProcess( mutex ) ;
        mp[2] = new MutexProcess( mutex ) ;

        // start MutexProcess
        mp[0].start() ;
        mp[1].start() ;
        mp[2].start() ;
    }
} // SEMADEMO
```

**The End.**